

Detecting loop-level parallelism

Tuesday, February 9, 2021 8:51 AM

Examples of no loop-carried dependencies

- o Jacobi iterations
- o SAXPY

for (i=0; i<N; i++)
y[i] = a * x[i] + y[i];

iter 0 iter 1
y[0] = ax[0] + y[0] y[1] = ax[1] + y[1] ...

Examples of loop-carried dependencies

- o Fibonacci series

How to have a compiler (such as gcc) automatically detect loop-carried dependencies?

- o Needed to safely perform loop unrolling
- o The GCD test

The programmer must get involved for harder loops --- rewrite the loop appropriately

- o Example: infinite series for approximating \pi

0 1 2 3 5 8 13 ...

fib[0] = 0

fib[1] = 1

#pragma omp parallel for

for (i=2; i<N; i++)
fib[i] = fib[i-1] + fib[i-2];

OpenMP does not check for loop carried dependencies

for (i=0; i<n; i++) {

A[i+b] = ...;

= A[c+i+d]

}

a, b : constant
c, d : constant

j, k : loop iterations
j < k
j ≠ k
0 ≤ j ≤ n; 0 ≤ k ≤ n

iterations

A[j+b] = ...;
= A[c+j+d];

A[k+b] = ...;
= A[c+k+d];

Race condition?

Iter j is writing to a location
Iter k is reading from a location

j+b == c+k+d *

Hazard

loop carried dependency

Test : greatest common divisor (GCD)

Test will pass if a loop-carried dependency exists

If a loop-carried dependency exists,
GCD(a, c) must divide (d-b)

Examples :

$$\text{for}(i=0; i<100; i++)$$

$$x[2i+3] = x[2i] + 5;$$

\downarrow
 a

\downarrow
 b

\downarrow
 c

\downarrow
 $d=0$

$$\text{GCD}(a, c) = \text{GCD}(2, 2) = 2$$

$$d - b = 0 - 3 = -3$$

2 does not divide -3
 \Rightarrow no loop carried dependency.

$$\text{for}(i=0; i<10; i++) \{$$

$$a[2i] = b[i];$$

$$c[i] = a[4i+1];$$

$$\}$$

$$a=2 \quad c=4$$

$$b=0 \quad d=1$$

$$\text{GCD}(a, c) = \text{GCD}(2, 4) = 2$$

$$d - b = 1 - 0 = 1$$

2 cannot divide 1

Sufficient condition

Programmer may need to get involved.

$$\pi = 4 \left[\frac{1}{1} - \frac{1}{3} + \frac{1}{5} - \frac{1}{7} + \dots \right] = 4 \sum_{k=0}^{\infty} \frac{(-1)^k}{2k+1}$$

Serial code:

plus_minus = 1;

sum = 0;

for(k=0; k<n; k++) {

sum += plus_minus (2k+1);

plus_minus = -plus_minus;

} approx_pi = 4 * sum;

\Rightarrow loop carried dependency

```

sum=0;
#pragma omp parallel for reduction(+:sum)
for (k=0; k<n; k++) {
    if (k%2 == 0)
        plus_minus = 1;
    else plus_minus = -1;
    sum += plus_minus / (2*k+1)
}
approx_pi = 4*sum;


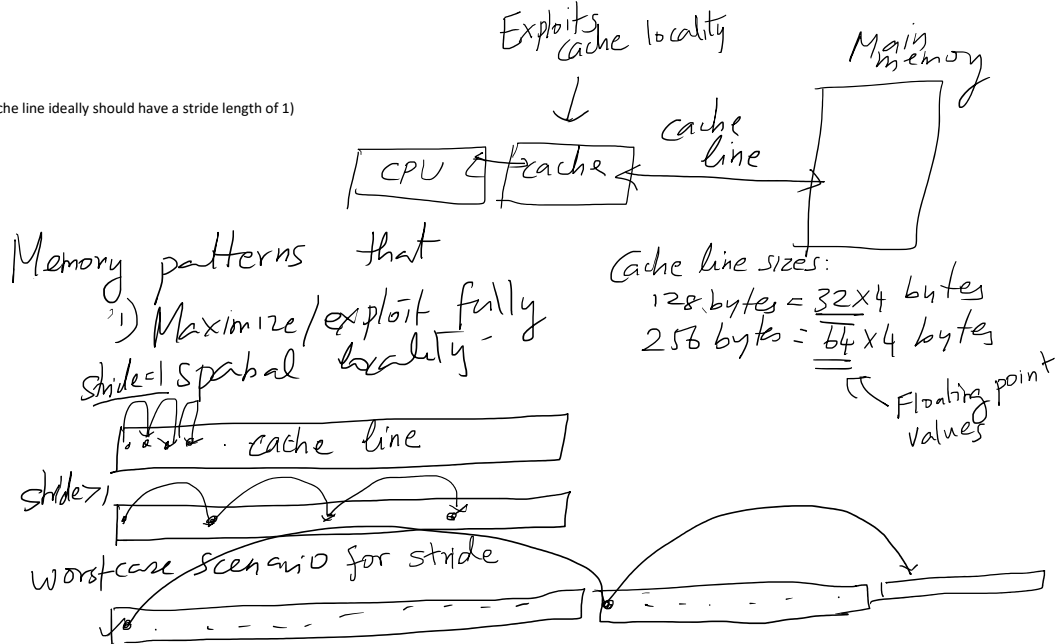
```

No loop carried dependency.

Loop optimizations

Tuesday, February 9, 2021 8:59 AM

- Approaches aim to:
 - o Maximize locality in the cache (memory accesses within a single cache line ideally should have a stride length of 1)
 - Optimize memory access patterns accordingly
 - o Reuse data brought into the cache
- Examples of loop optimization:
 - o Loop fusion
 - o Loop fission
 - o Loop interchange
 - o Loop tiling
 - Tiled matrix transpose
 - Tiles matrix multiplication



1D array in row-major form:

row 0

row 1

...

row n-1

$$A[i][j] = A[i \times \text{num_cols} + j]$$

num_cols : # of columns (or) width of matrix

Matrix transpose: assume $n \times n$ matrix

```
for (i=0; i<n; i++)
    for (j=0; j<n; j++)
        B[i][j] = A[j][i];
```

Memory access along columns of A



- ↳ Very large striding of the 1D layout in memory
- ↳ Very large number of cache misses

Solution: maximize the reuse of data brought into cache

→ Operate on smaller-sized tiles, one at a time (see loop-tiling-c)

Matrix multiplication:

$P = M \times N$ Assume $n \times n$ matrices



p_{ij} = dot product of
 $\left(\begin{matrix} i^{\text{th}} \text{ row} \\ \text{of } M \end{matrix} \right)$ and $\left(\begin{matrix} j^{\text{th}} \text{ column} \\ \text{of } N \end{matrix} \right)$

Performance problem: column access for matrix N
 \rightarrow Each access to a column element results in a cache miss (due to large stride)

Solution: tiled/blocked matrix multiplication.

Introduction to CUDA

Tuesday, February 9, 2021 9:03 AM

- Compute Unified Device Architecture (CUDA) is a parallel computing platform and programming model developed by NVIDIA for general computing on graphical processing units (GPUs)
- Enables General Purpose Programming on the GPU (GPGPU)
- Contains SPMD extensions for data-parallel programming on the GPU
 - o SPMD: Single Program Multiple Data
 - o STMD: Single Thread Multiple Data

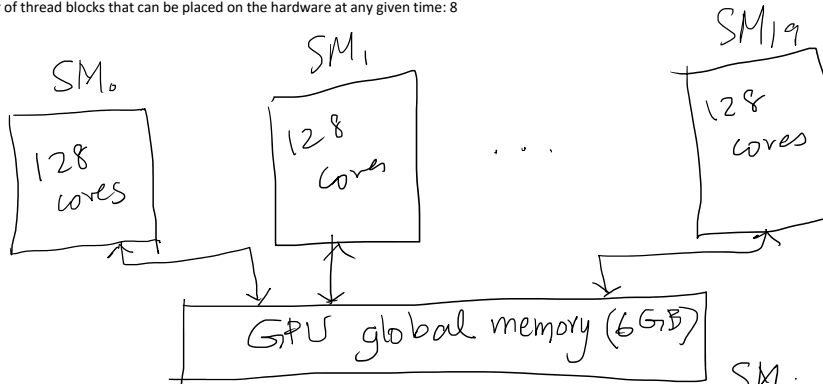
GPU microarchitecture

Tuesday, February 9, 2021 9:05 AM

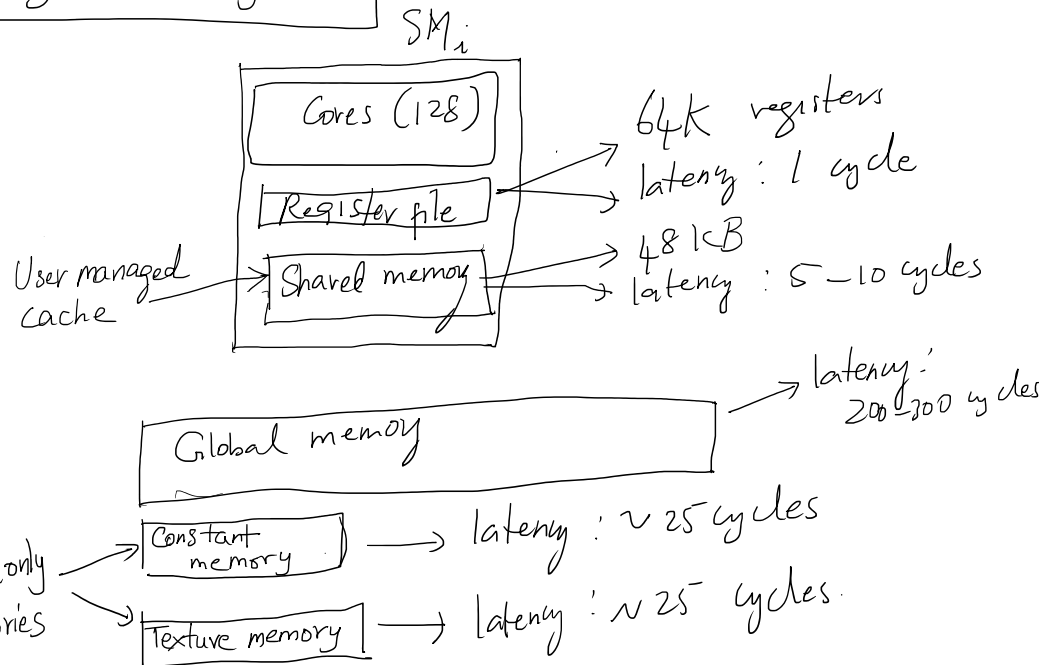
- We will be using the GTX 1080 on xunil-05.coe.drexel.edu
- Feel free to debug and test code on your local GPU (if you have one on your machine)
- The GTX 1080 comprises 2560 cores
 - o 20 streaming multiprocessors (SMs), each with 128 cores
 - o 64K registers per SM
 - o 48 KB of shard memory/cache per SM
 - o 6 GB of global memory
 - o 64 KB of constant memory
- Maximum number of threads that can be executed by the hardware at any given time: 2048 threads
- Maximum number of thread blocks that can be placed on the hardware at any given time: 8

Cores are organized in hierarchical fashion

- Operates at $\sim 1.6 \text{ GHz}$
- Supports single precision and double precision operations
- Supports branch statements



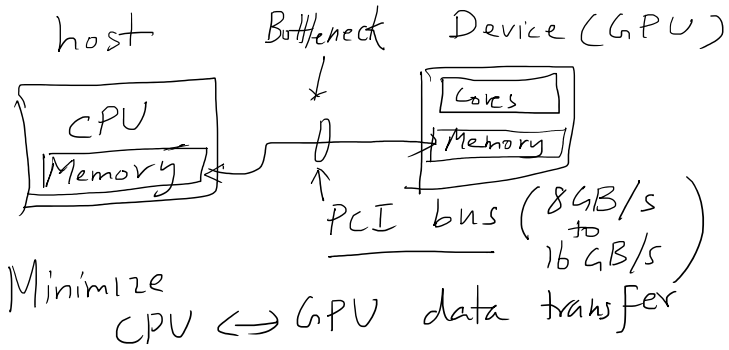
GPU memory hierarchy:



The CUDA/C execution model

Tuesday, February 9, 2021 9:09 AM

- Integrated host and GPU program
- Serial or modestly parallel parts in host C code
- Highly parallel parts in GPU SPMD (single program multiple data) kernel C code
- CPU and GPU maintain separate memory maps
 - o Use `cudaMemcpy()` to transfer data between CPU and GPU
- Memory allocation on the GPU is persistent across multiple kernel calls
 - o `cudaMalloc()` allocates memory on the GPU
 - o `cudaFree()` frees previously allocated memory



Host side

1. Allocate memory on GPU
 2. Transfer data to GPU (CPU → GPU)
 3. Setup execution grid
 4. Launch kernel
 5. Transfer data to CPU (GPU → CPU)
6. Free GPU global memory

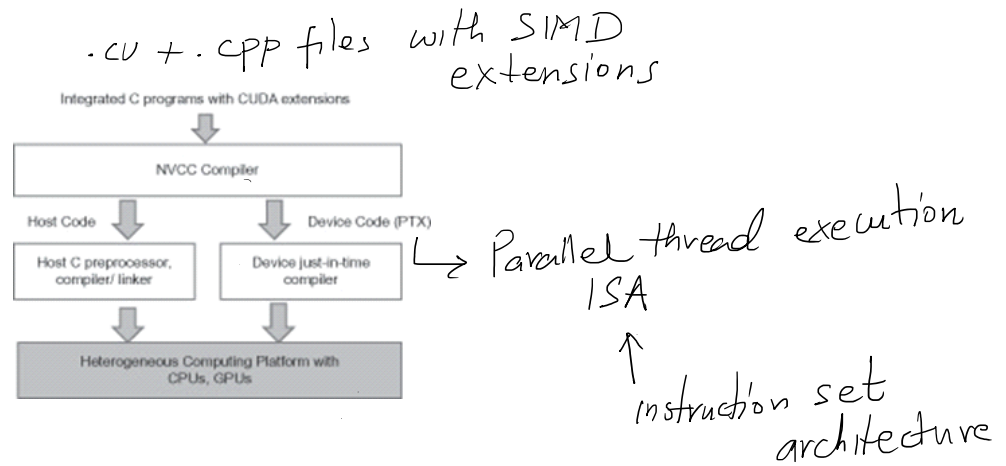
Device side

- kernel

GPU Global memory is persistent across multiple kernel calls.

Building the CUDA executable

Tuesday, February 9, 2021 9:16 AM





CUDA threading model

Tuesday, February 9, 2021 9:18 AM

- Hierarchical thread organization
 - o Threads
 - o Thread blocks
 - o Execution grid
- CUDA provides intrinsic variables to help locate each thread within the execution grid
 - o Needed for thread-to-data mapping for your application (think thread ID)
- More on thread blocks
 - o Upper bound on thread-block size (CUDA version/GPU architecture dependent)
 - Maximum number of threads/thread block = 1024 on the 1080 GTX
 - Examples: (1024, 1, 1), (1, 1024, 1), (32, 32, 1)
 - o Thread block must be scheduled in its entirety on a streaming multiprocessor; cannot be split up over multiple streaming multiprocessors
 - o Maximum number of blocks that can be scheduled on a SM at any time: 8
 - o Threads can use shared memory and synchronization mechanisms (barriers) to coordinate with other threads within a single thread block
 - o No synchronization is available or assumed between thread blocks (that is, thread blocks can finish execution in any order)
- Upper bound on execution-grid size
 - o Maximum number of thread blocks in the X direction: 2^{16}
 - o Maximum number of thread blocks in the Y direction: 2^{16}

Hierarchical organization



Thread blocks

1D thread block:

Intrinsic variables:


$$tid = threadIdx.x \cdot x;$$

max of threads/block: 1024

$(1024, 1, 1)$

2D thread block:

threadIdx.x
threadIdx.y
blockDim.x
blockDim.y



$(threadIdx.x \cdot x, threadIdx.x \cdot y)$


Dimensions of thread block:

blockDim.x
blockDim.y

(# of threads along X)
(# of threads along Y)

Execution grid:

1D grid:




blockIdx.x
blockIdx.y

locates thread within grid.

$$tid = blockIdx.x \cdot blockDim.x + threadIdx.x;$$

2D grid :



$$\text{row} = \text{blockIdx.y} * \text{blockDim.y} + \text{threadIdx.y};$$

$$\text{col} = \text{blockIdx.x} * \text{blockDim.x} + \text{threadIdx.x};$$


gridDim.x (# threads in x direction)
gridDim.y (# threads in y direction)

Example: vector addition

Tuesday, February 9, 2021 11:35 AM

A, B : vectors

$$C = A + B$$



- 1) Strategy: each output element in C is computed by a thread.
- 2) If I am a thread, what input elements do I need to calculate my output element

Code example (single thread block)

Sunday, May 03, 2020 11:16 AM

thread block : (1024, 1, 1) \xrightarrow{x} $\{S_0, S_1, \dots, S_{1023}\}$
grid : (1, 1) \xrightarrow{y} $\{TB\}$


kernel code :

$tid = threadIdx.x;$
 $C[tid] = A[tid] + B[tid];$

A : $\{ \dots \}$ 1023
+ B : $\{ \dots \}$
|| C : $\{S_0, S_1, \dots, S_{1023}\}$
0 1 1023

Arrays A and B contain exactly 1024 elements.

A and B are vectors of arbitrary length n



Execution grid : thread block : $(1024, 1, 1)$
 # thread blocks : $\left\lceil \frac{n}{\text{thread block size}} \right\rceil = \left\lceil \frac{n}{1024} \right\rceil$
 grid : $(\# \text{ thread blocks}, 1)$

kernel code :

```

tid = blockIdx.x * blockDim.x + threadIdx.x;
if (tid < n) {
    C[tid] = A[tid] + B[tid];
}

```

Performance issue : Creation of large number of thread blocks \Rightarrow more work for the CUDA runtime to move thread blocks in and out of the hardware

Version 2 :

- ↖ fixed
- Create limited number of thread blocks
- Have each thread calculate multiple output elements in C

How many thread blocks to create?

"Back of envelop" calculation:

- Each SM can accommodate 2048 threads at any time.
- 2 thread blocks of dimension $(1024, 1, 1)$ per SM

. So, # of thread blocks = $2 \times 20 = 40$


Execution grid:

thread block: (1024, 1, 1)

#thread blocks: 40

grid: (#thread blocks, 1) = (40, 1)

We use the concept of striding to calculate output elements.



$$\text{Stride length} = \text{gridDim.x} * \text{blockDim.x};$$

total number of threads in the grid.

Kernel code:

$\text{tid} = \text{blockIdx.x} * \text{blockDim.x} + \text{threadIdx.x};$

$\text{stride} = \text{gridDim.x} * \text{blockDim.x};$

while (tid < n) {

$\text{c}[\text{tid}] = \text{A}[\text{tid}] + \text{B}[\text{tid}];$ // Calculate current element

$\text{tid} += \text{stride};$ // Move on to next element

}