Parallel reduction

Tuesday, February 23, 2021 8:22 AM

- Code example: vector_reduction_small

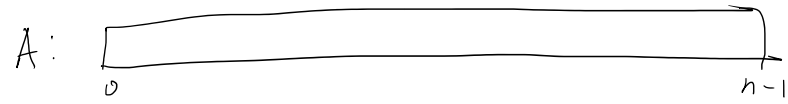

$$S = \sum_i a_i$$

V1

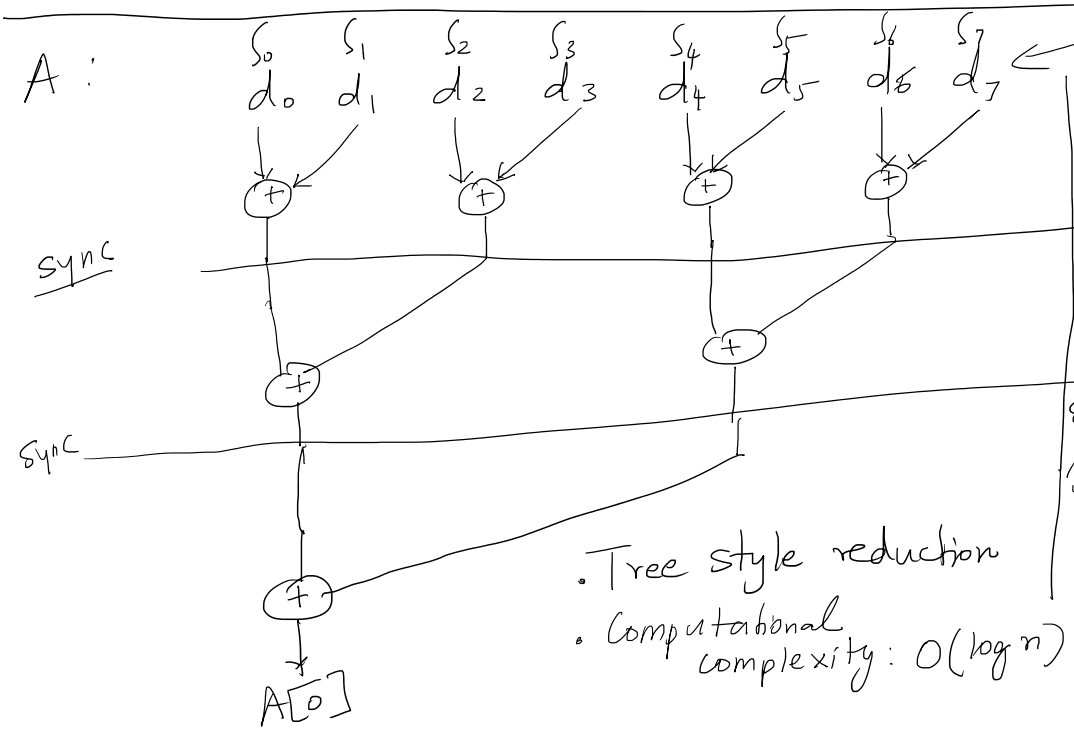

loaded elements from global to shared memory.

```
stride = 1
if (tid/2*stride == 0)
    A[tid] += A[tid+stride]
stride = 2 * stride = 2
if (tid/2*stride == 0)
    A[tid] += A[tid+stride]
stride = 2 * stride = 4
if (tid/2*stride == 0)
    A[tid] += A[tid+stride]
```

- Tree style reduction
- Computational complexity: $O(\log n)$

Performance issue:

- Branch divergence within warp

V2 Tree style reduction

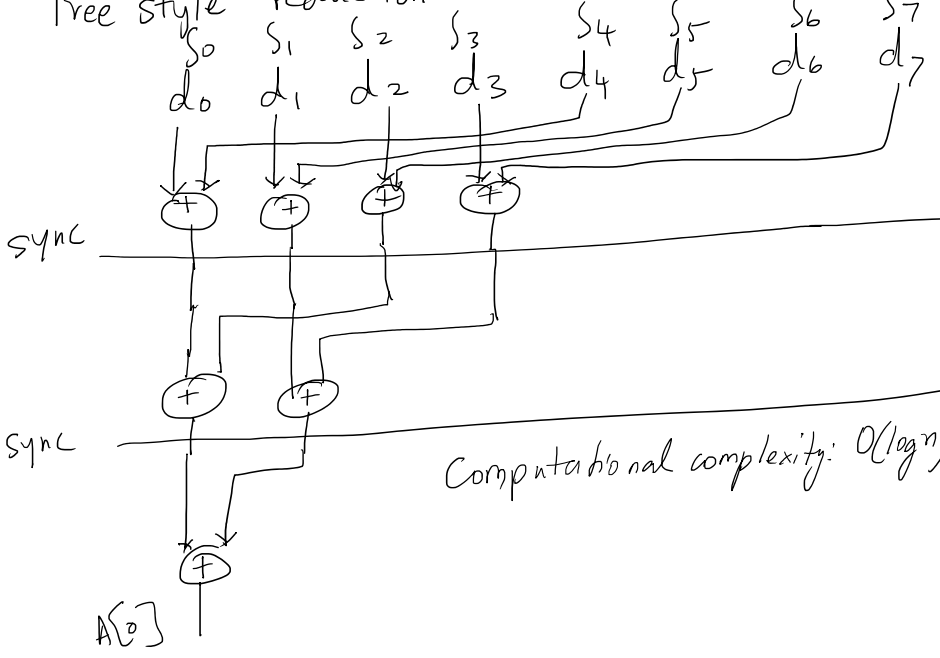

Computational complexity: $O(\log n)$

```
stride = 8/2 = 4
if (tid < stride)
    A[tid] += A[tid+stride]
stride = stride/2 = 4/2 = 2
if (tid < stride)
    A[tid] += A[tid+stride]
stride = stride/2 = 2/2 = 1
if (tid < stride)
    A[tid] += A[tid+stride]
```

Version V2 shows reduced branch divergence:

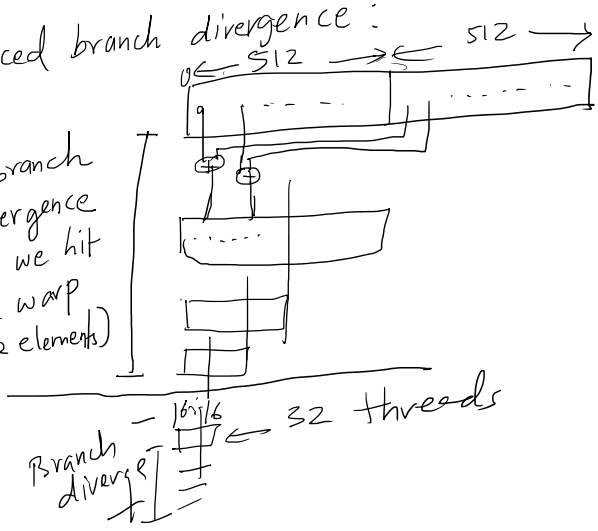

- # of elements must be power of 2 (since it is tree style)
- What if not power of 2?
  - Pad out shared memory area with zeros to nearest power of 2

Reducing large arrays:

A:

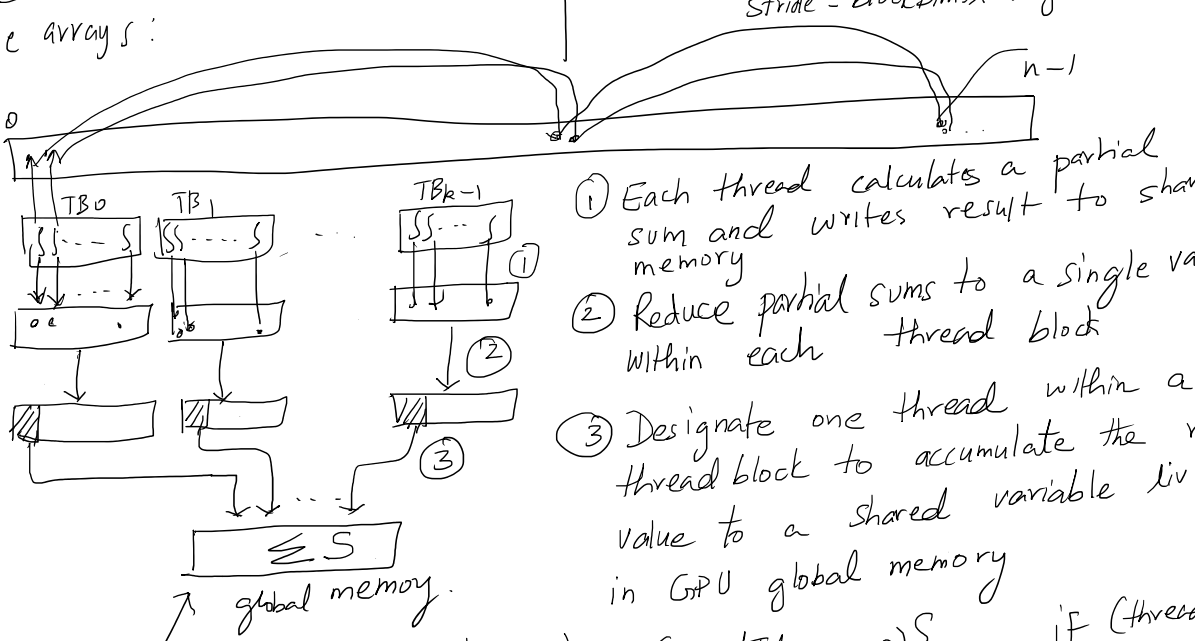

① Each thread calculates a partial sum and writes result to shared memory

② Reduce partial sums to a single value within each thread block

③ Designate one thread within a thread block to accumulate the reduced value to a shared variable living in GPU global memory

```
if (threadIdx.x == 0) {
    atomicCAS(mutex);
    // accumulate into S
    mutex = 0;
}
```

(or) `if (threadIdx.x == 0) atomicAdd(S)`

Shared variable must be protected via a mutex or use atomic operations supported by NVIDIA's ISA.

Parallel scan

Tuesday, February 23, 2021 8:22 AM

- Code example: scan

$\oplus$: scan operator
ex: addition

I : identity element

Scan:

input : $[a_0\ a_1 \dots a_{n-1}]$

Inclusive scan:

output : $[a_0,\ a_0 \oplus a_1\ \ a_0 \oplus a_1 \oplus a_2,\ \dots\ a_0 \oplus a_1 \oplus \dots \oplus a_{n-1}]$

Exclusive scan:

output : $[I\ \ a_0\ \ a_0 \oplus a_1\ \ a_0 \oplus a_1 \oplus a_2 \dots a_0 \oplus a_1 \oplus \dots \oplus a_{n-2}]$

Example: scan operator = +

Input : [3 1 7 0 4 1 6 3]
IScan : [3 4 11 11 15 16 22 25]
EScan : [0 3 4 11 11 15 16 22]

Serial code:

Escan:

```
scan (input, output, n) {
    output[0] = 0;
    for (i=1; i<n; i++)
        output[i] = output[i-1] + input[i];
}
```

Parallel prefix scan:

Inclusive scan

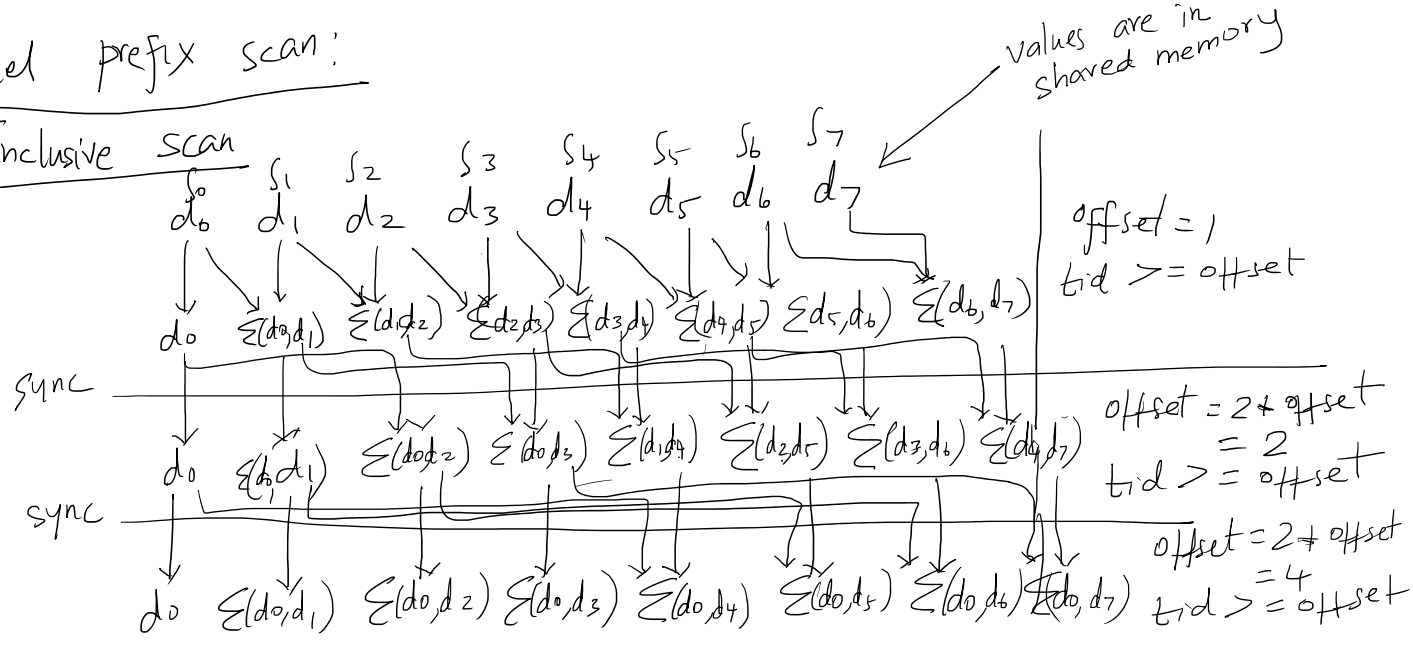

Use ping-pong buffers to store values between steps to prevent read after write hazards

Recap: atomics on the GPU

Thursday, February 25, 2021 8:14 AM

- Use *atomicCAS()* to achieve mutual exclusion to a critical section of code
  - Do not use *atomicCAS()* as a synchronization mechanism between threads belonging to the same thread warp --> may lead to deadlocks
  - Use *atomicCAS()* as synchronization mechanism between thread blocks
    - Ideally: designated one thread in each thread block, say threadidx.x = 0, to be the writer to the critical section:

```
If (threadIdx.x == 0) {
        Lock()
        /* Critical section code */
        Release()
}
```

- For operations such as addition, subtraction, etc., it is preferrable to use the atomic variants of these operations for better performance:
  - *atomicAdd(), atomicSub()*, …
  - See: https://docs.nvidia.com/cuda/cuda-c-programming-guide/index.html#atomic-functions

Recap: parallel scan

Thursday, February 25, 2021 8:23 AM

Scan : +

Input : [3 1 7 0 4 1 6 3]

iScan : [3 4 11 11 15 16 22 25]

eScan : [0 3 4 11 11 15 16 22]

eScan:

```
output[0] = 0
for (i = 1; i < n; i++)
  output[i] = output[i-1] + input[i-1];
```

↑ loop carried dependency

Parallel prefix scan:

Inclusive scan:

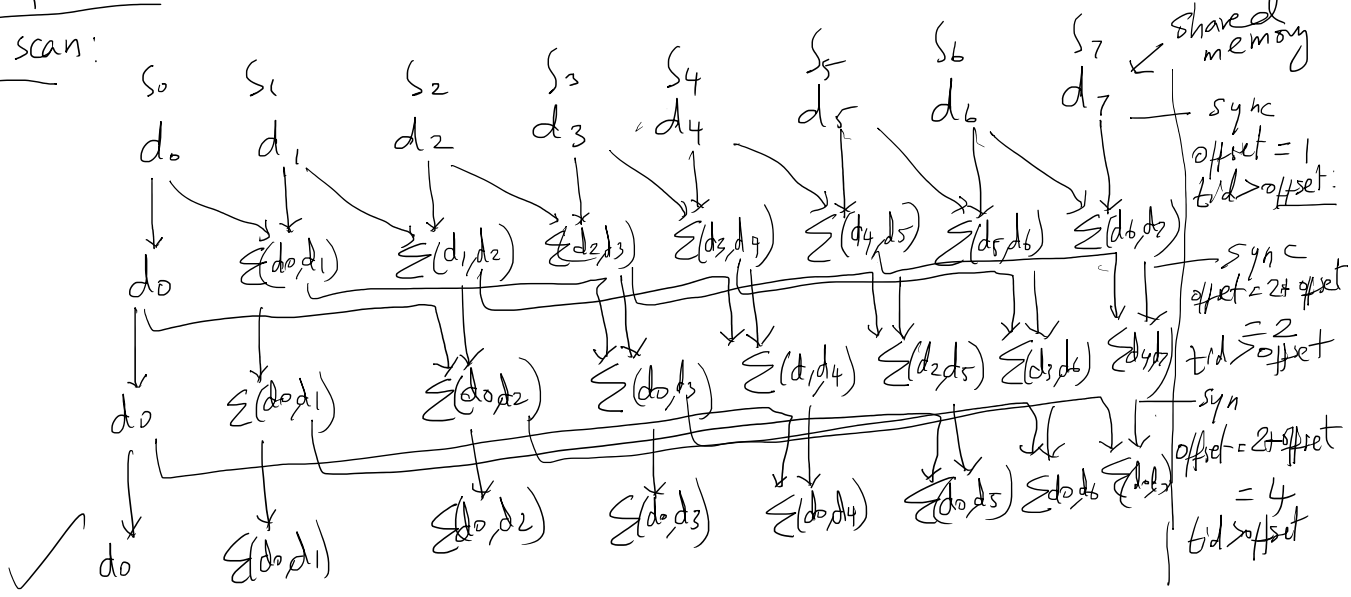

Exclusive Scan:

- Use same technique as above
- Threads load data into shared memory as

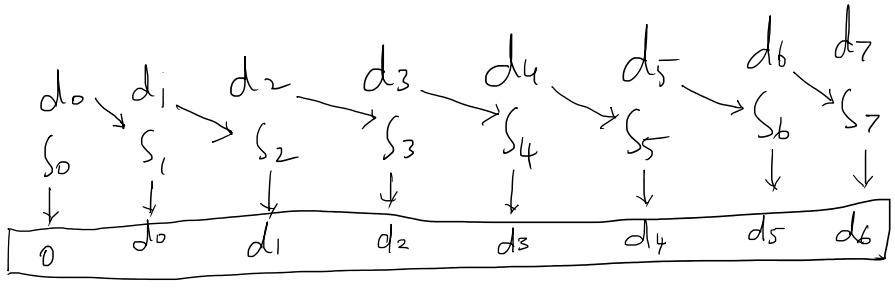

Stream compaction

Thursday, February 25, 2021 8:23 AM

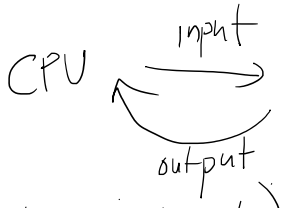

(filtered input)

filter:
only +ve values
> 0
<= 0 ×

size(output) ≤ size(input)

Input: 1 -1 3 4 -6 5 -8 10
Output: 1 3 4 5 10

	S_0	S_1	S_2	S_3	S_4	S_5	S_6	S_7
input	1	-1	3	4	-6	5	-8	10
flag	1	0	1	1	0	1	0	1
escan(flag)	0	1	1	2	3	3	4	4

↑ locations in the output array for each thread.

Output: 1 3 4 5 10
0 1 2 3 4

elements = 5

Counting sort

Thursday, February 25, 2021 8:54 AM

Counting sort : $D: \{d_0, d_1, \ldots d_{n-1}\}$ $d_i \in$ range (integers) (for example, [0, 255])

- Non-comparison based
- Complexity : O(n)

input : 8 5 1 3 7 8 6 5 3 8

bin :

S_0	S_1	S_2	S_3	S_4	S_5	S_6	S_7	S_8
bin 0	bin 1	bin 2	bin 3	bin 4	bin 5	bin 6	bin 7	bin 8
0	1	0	2	0	2	1	1	3
0	0	1	1	3	3	5	6	7

create histogram

generate starting address for each thread in sorted array

escan(bin)

output

1	3	3	5	5	6	7	8	8	8
0	1	2	3	4	5	6	7	8	9

Histogram generation

Tuesday, February 23, 2021 8:22 AM

- Code example: histogram_generation

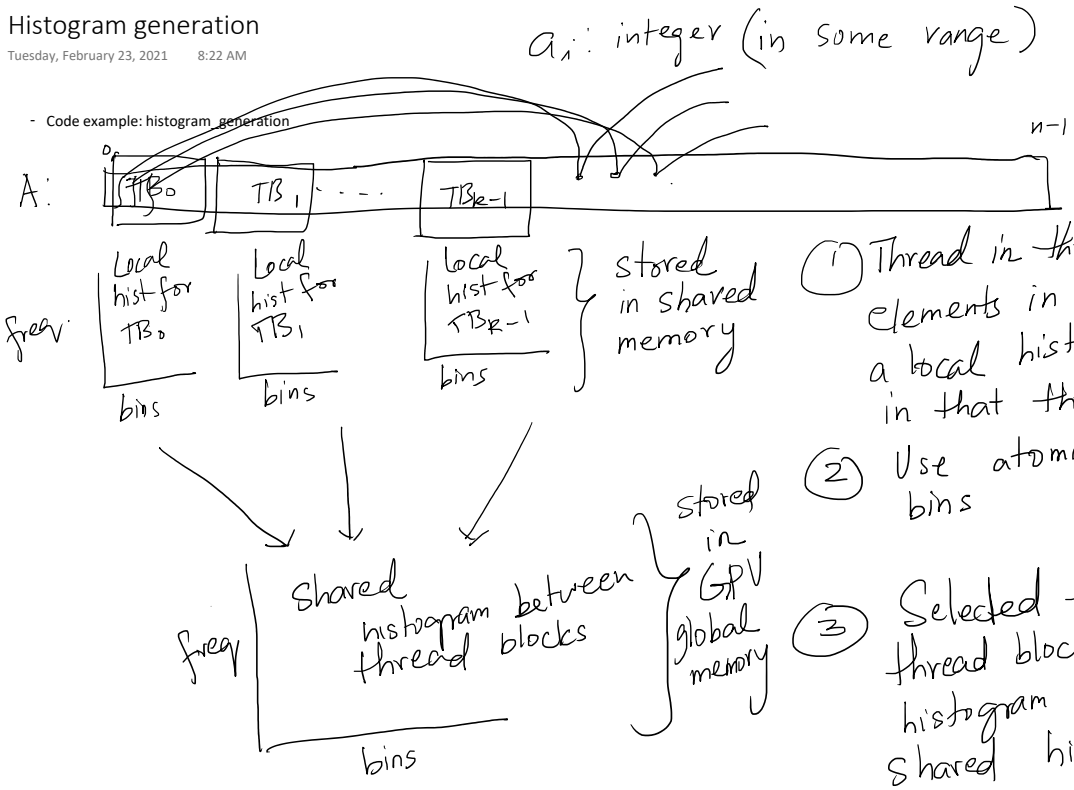

(1) Thread in thread block strides across elements in A and bins each in a local histogram shared by threads in that thread block

(2) Use atomics when incrementing bins

(3) Selected threads within each thread block accumulate local histogram bin values into the shared histogram in GPU global memory

(4) Use atomics when accumulating bin values in the shared histogram